

Session 1 - R Basics

Jongbin Jung

January 9-10, 2016

A Question

You are given three sticks, each of a random length between 0 and 1.

What's the probability you can make a triangle?

- ▶ The answer is $1/2$
- ▶ By the end of this session, we'll confirm this with a simulation in R

Dependencies

- ▶ Latest version ($\geq 3.1.2$) of R
(*free* from <https://www.r-project.org/>)
- ▶ Latest version of Rstudio (also *free* from <https://www.rstudio.com/>)
- ▶ A bunch of *free* packages

```
install.packages('dplyr')  
install.packages('ggplot2')
```

Rstudio Basics

The screenshot displays the RStudio IDE interface. The top menu bar includes File, Edit, Code, View, Plots, Session, Build, Debug, Tools, and Help. The toolbar below the menu contains icons for file operations and execution. The main editor window shows a script named 'notes.Rmd' with R Markdown code. The code includes a YAML header for a Beamer presentation, dependencies for R 3.1.2, Rstudio, and various packages (dplyr, ggplot2), and a section for Rstudio basics. The console window at the bottom shows the output of the code, including summary statistics for a variable 'speed' and a summary of the 'cars' dataset. The file explorer on the right shows the project structure, including folders for 'gignore', '.Rhistory', 'r-tutorial.Rproj', 'README.md', 'refs', 'Session 1', and 'img'.

```
6 beamer_presentation:
7   highlight: zenburn
8 ---
9
10 ## Dependencies
11
12 - Latest version ( $\$ge\$`3.1.2`$ ) of 'R'
13   (*free* from <https://www.r-project.org/>)
14 - Latest version of 'Rstudio'
15   (also *free* from <https://www.rstudio.com/>)
16 - A bunch of *free* packages
17
18 ```{r, eval=FALSE}
19 install.packages('dplyr')
20 install.packages('ggplot2')
21 ```
22
23 ## 'Rstudio' Basics
24
25 - 'Ctrl' + 'Enter': Execute current line of code
26
26.11 (Top Level) >
```

Environment: Global Environment (empty)

Files: Home / repos / r-tutorial

Name	Size	Modified
..		
gignore	213 B	Dec 18, 2015, 11:38 AM
.Rhistory	0 B	Dec 18, 2015, 11:38 AM
r-tutorial.Rproj	235 B	Dec 18, 2015, 11:38 AM
README.md	24 B	Dec 17, 2015, 8:41 PM
refs		
Session 1		
img		

Console: ~/repos/r-tutorial/

```
list OD: 12.0 1st Qu.: 28.00
Median: 15.0 Median: 36.00
Mean: 15.4 Mean: 42.98
3rd Qu.: 19.0 3rd Qu.: 56.00
Max.: 25.0 Max.: 120.00
> summary(cars)
      speed      dist
Min.   : 4.0   Min.   : 2.00
1st Qu.:12.0   1st Qu.: 26.00
Median :15.0   Median : 36.00
Mean   :15.4   Mean   : 42.98
3rd Qu.:19.0   3rd Qu.: 56.00
Max.   :25.0   Max.   :120.00
> |
```

Rstudio Basics

- ▶ `Ctrl + #`: Focus on panel #
- ▶ `Ctrl + Enter`: Execute selection
- ▶ `Ctrl + Shift + C`: Comment/Uncomment selection
- ▶ Many more (if you're willing to explore)

R Basics

R Basics: Working Directory

- ▶ Working directory (wd) is where your R session will load/save files
- ▶ To see where your current working directory is, run

```
getwd()
```

- ▶ To set the working directory to desired path, run

```
setwd("path")
```

- ▶ Note that ~ is replaced with your HOME directory, e.g. C:\Users\Username\ in windows
- ▶ Use forward slashes (/), even on Windows!

```
setwd("~/Documents/R")
```

R Basics: Math Operations

- ▶ Simple math operations

```
3+11 # add stuff  
3-11 # subtract stuff  
3/11 # divide stuff  
3*11 # multiply stuff  
2^10 # raise to powers
```


R Basics: Assignments

- ▶ Convention for assigning values to variables is `<-`
- ▶ Direction of arrow indicates direction of assignment

```
A <- 12
A # 12
A + 3 -> B
B # 15
24 -> A
A # 24
```

- ▶ The equal sign (`=`) also works, but only for assignment to the left, e.g.

```
A = 12 # good
12 = A # BAD
```

R Basics: Strings

- ▶ A String variable can be declared in either double quotes("") or single quotes ('')

```
str <- "This is a valid string"  
str
```

```
## [1] "This is a valid string"
```

```
str <- 'and so is this'  
str
```

```
## [1] "and so is this"
```

R Basics: Re-Assignments

- ▶ A variable can be re-assigned to anything

```
x <- 860306 # first x is assigned a number  
x
```

```
## [1] 860306
```

```
x <- 'This is a variable!'  
x # Now it is a string
```

```
## [1] "This is a variable!"
```

Vectors

Vectors: c()

- Vectors in R are created by concatenating a series of elements

```
X <- c(1,2,3)
X # vector of numbers (1, 2, 3)
```

```
## [1] 1 2 3
```

```
Y <- c('this', 'that', 'those')
Y # this is a vector of Strings
```

```
## [1] "this" "that" "those"
```

Vectors: seq()

- ▶ Create a vector from a sequence with `seq(from, to, by=1)`

```
seq(1, 10)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
seq(1, 10, 2)
```

```
## [1] 1 3 5 7 9
```

- ▶ Use short-hand `from:to` if you're incrementing by one

```
1:10
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

Vectors: rep()

- Use rep() to repeat values

```
rep(13, 4)
```

```
## [1] 13 13 13 13
```

```
rep('Yes!', 3)
```

```
## [1] "Yes!" "Yes!" "Yes!"
```

```
rep(c('Sat.', 'Sun.'), 2)
```

```
## [1] "Sat." "Sun." "Sat." "Sun."
```

Vectors: `rdist()`

- ▶ Generate vector of n samples from a specified distribution

```
runif(n = 10)  # 10 samples from Unif(0, 1)
rnorm(n = 10)  # 10 samples from Norm(0, 1)
rpois(n = 10)  # 10 samples from Poisson(1)
rexp(n = 10)   # 10 samples from Exp(1)
```

- ▶ Distribution parameters can be specified as arguments, e.g.

```
# 100 samples from a Norm(20, 5) distribution
rnorm(n = 100, mean = 20, sd = 5)
```

- ▶ Read documentation for available distributions

```
?Distributions
```


Vectors: Indexing

- ▶ Use square braces ([]) to index a vector (base 1)

```
X <- c(10, 11, 12, 13)  
X[1]
```

```
## [1] 10
```

```
X[4]
```

```
## [1] 13
```

```
X[5]
```

```
## [1] NA
```

Vectors: Indexing (cont'd)

- ▶ Negative indexing is used to exclude elements

```
X[-1]
```

```
## [1] 11 12 13
```

- ▶ Index multiple objects by indexing with a vector

```
ind <- c(2, 4)  
X[ind]
```

```
## [1] 11 13
```

Vectors: Re-assignment with Indices

- ▶ Replace elements by re-assigning with index

```
X[1] <- 101  
X
```

```
## [1] 101 11 12 13
```

- ▶ Replace multiple elements as well

```
X[2:3] <- c(22, 33)  
X
```

```
## [1] 101 22 33 13
```

Vectors: Add Elements by Index

- Add new elements to a vector by assigning

```
X[5]
```

```
## [1] NA
```

```
X[5] <- 555  
X
```

```
## [1] 101 22 33 13 555
```

Vectors: Advanced Indexing

- ▶ Vectors can be indexed by a binary vector (TRUE/FALSE) of equal length
- ▶ i.e., you can index vectors by a specified condition, e.g.,

```
X <- 1:100  
# create a binary vector with the same length of X  
# where the element is TRUE if the element of X  
# in the corresponding position satisfies condition  
ind <- X > 95  
tail(ind) # take a peek at the last few entries
```

```
## [1] FALSE TRUE TRUE TRUE TRUE TRUE
```

```
X[ind]
```

```
## [1] 96 97 98 99 100
```

Matrices

Creating Matrices

- ▶ A matrix is created from a vector, using `matrix()`, e.g.

```
X <- c(1:12)
# syntax: matrix(vector, # of rows, # of columns)
A <- matrix(X, 3, 4)
A
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    4    7   10
## [2,]    2    5    8   11
## [3,]    3    6    9   12
```

- ▶ Notice that the matrix is created column-first

Matrix Indexing

- ▶ Similar to vectors, a matrix can be indexed with square braces, with the syntax `[row #, col #]`, e.g.
- ▶ Leaving an entry empty will result in the full row/column

```
A[3,2] # third row, second column
```

```
## [1] 6
```

```
A[2,] # entire second row
```

```
## [1] 2 5 8 11
```

```
A[,4] # entire fourth column
```

```
## [1] 10 11 12
```


Vector/Matrix Operations

Vector Operations

```
X = c(1:4)
t(X)  # transpose (column) vector X to row vector
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    2    3    4
```

```
X + X  # element-wise summation
```

```
## [1] 2 4 6 8
```

```
X - X  # element-wise subtraction
```

```
## [1] 0 0 0 0
```

Vector Operations (cont'd)

```
X^3    # element-wise exponentiation
```

```
## [1]  1  8 27 64
```

```
X * X  # element-wise multiplication
```

```
## [1]  1  4  9 16
```

```
X %*% X  # dot (inner) product
```

```
##      [,1]
```

```
## [1,]    30
```

Matrix Operations

```
A = matrix(1:4, 2, 2) # create 2x2 matrix  
t(A) # transpose (column) vector A to row vector
```

```
##      [,1] [,2]  
## [1,]    1    2  
## [2,]    3    4
```

```
A + A # element-wise summation
```

```
##      [,1] [,2]  
## [1,]    2    6  
## [2,]    4    8
```

```
A - A # element-wise subtraction
```

```
##      [,1] [,2]  
## [1,]    0    0  
## [2,]    0    0
```

Matrix Operations (cont'd)

```
A^3 # element-wise exponentiation
```

```
##      [,1] [,2]  
## [1,]    1  27  
## [2,]    8  64
```

```
A * A # element-wise multiplication
```

```
##      [,1] [,2]  
## [1,]    1    9  
## [2,]    4   16
```

```
A %*% A # dot (inner) product
```

```
##      [,1] [,2]  
## [1,]    7   15  
## [2,]   10   22
```

Matrix Operations: Warning

- Dimensions must make sense!

```
A <- matrix(1:6, 2, 3)
B <- matrix(1:6, 3, 2)
A %*% B # 2x3 times 3x2: OK
```

```
##      [,1] [,2]
## [1,]   22  49
## [2,]   28  64
```

```
A %*% A # 2x3 times 2x3: Nope
```

```
## Error in A %*% A: non-conformable arguments
```

Vector/Matrix Comparisons

- ▶ Comparisons are all done element-wise

```
c(1, 2, 3) == c(1, 2, 4)
```

```
## [1] TRUE TRUE FALSE
```

```
c(1, 2, 3) < c(1, 2, 4)
```

```
## [1] FALSE FALSE TRUE
```

```
c(1, 2, 3) >= c(1, 2, 4)
```

```
## [1] TRUE TRUE FALSE
```

- ▶ Note the double equal sign for comparing equality (one would be assignment!)

Helpful Vector Functions

- If possible, avoid loops by operating over the Vector/Matrix as a whole

```
mean(X)           # mean
sd(X)             # standard deviation
var(X)            # variance
max(X)            # maximum
min(X)            # minimum
median(X)         # median
sum(X)            # sum
prod(X)           # product
quantile(X,probs=0.5) # quantile for specified probs
length(X)         # length of the vector
range(X)          # range
```


Helpful Matrix Functions

- If possible, avoid loops by operating over the Vector/Matrix as a whole

```
rowSums(A)  # Row sums
colSums(A)  # Column sums
rowMeans(A) # Row means
colMeans(A) # Columns means
diag(A)     # Diagonal of a matrix
solve(A)    # Inverse of a matrix
cov(A)      # Variance covariance matrix
cor(A)      # Correlation matrix
```

Functions

Some more built-in functions

- ▶ We've already seen many built-in functions, but here are some more!

```
log(X)    # element-wise log
exp(X)    # element-wise exponential
sqrt(X)   # element-wise square root
```

Functions for Strings

```
# concatenate two (or more) strings  
paste('one plus one equals', 1+1, '!')
```

```
## [1] "one plus one equals 2 !"
```

```
# specify a separator  
paste('one plus one', 1+1, sep='=')
```

```
## [1] "one plus one=2"
```

```
# if you're into C-style formatting ...  
sprintf('one plus one = %d', 1+1)
```

```
## [1] "one plus one = 2"
```

Functions for Strings (cont'd)

- ▶ Often, we want to concatenate strings with no spaces (e.g., when constructing filenames/paths in run-time)

```
# short-hand for concatenation w/o spaces  
filename = 'some_file_name'  
paste0('path/to/', filename, '.csv')
```

```
## [1] "path/to/some_file_name.csv"
```

```
paste0(getwd(), '/', filename, '.csv')
```

```
## [1] "/home/jongbin/repos/r-tutorial/1-intro/some_file_name.csv"
```

Functions for Strings (cont'd)

- ▶ To enforce upper/lower cases

```
s <- 'SoMe CraZY STRING'  
tolower(s)
```

```
## [1] "some crazy string"
```

```
toupper(s)
```

```
## [1] "SOME CRAZY STRING"
```

Generic Functions

- Some functions for exploring objects

```
obj <- 1:100  
head(obj, n=5)  # display first n rows of obj
```

```
## [1] 1 2 3 4 5
```

```
tail(obj, n=5)  # display last n rows of obj
```

```
## [1] 96 97 98 99 100
```

Generic Functions (cont'd)

```
str(obj) # display structure of obj
```

```
## int [1:100] 1 2 3 4 5 6 7 8 9 10 ...
```

```
summary(obj) # display summary of obj
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.00   25.75   50.50   50.50   75.25   100.00
```


Control Flow

► if statements

```
if (condition) {  
    # stuff to do when condition is TRUE  
} else if(other_condition) { # (OPTIONAL)  
    # stuff to do if other_condition is TRUE  
} else { # (OPTIONAL)  
    # stuff to do if all conditions are TRUE  
}
```

Loops

- ▶ for statements

```
for (ind in sequence/set) {  
    # iterate over sequence or elements of a set  
    # do stuff  
}
```

- ▶ while statements

```
while (condition) {  
    # stuff to do while the condition is TRUE  
    # the condition must become FALSE at some point!  
}
```

Loops: Example

```
for (i in 1:3) {  
  print(paste('iteration', i))  
}
```

```
## [1] "iteration 1"  
## [1] "iteration 2"  
## [1] "iteration 3"
```

```
while (i >= 0) {  
  print(paste('de-iteration', i))  
  i <- i - 1 # beware of infinite loops!  
}
```

```
## [1] "de-iteration 3"  
## [1] "de-iteration 2"  
## [1] "de-iteration 1"  
## [1] "de-iteration 0"
```

User Defined Functions

- ▶ Write your own functions in the form

```
name_of_function <- function(arguments) {  
  # do some stuff with arguments  
  return(result)  
}
```

- ▶ You can use your functions like any other function, e.g.,

```
name_of_function(arguments) # gives you the 'result'
```

User Defined Functions: Example

- Write a function that will take a vector in \mathbb{R}^3 and tell you if you can make a triangle or not

```
is_good <- function(vec) {  
  if (length(vec) != 3) {  
    # it's always a good idea to make sure your  
    # function gets what it expects to get  
    stop('is_good requires a vector of length 3')  
  }  
  for (i in 1:3) {  
    if (vec[i] > sum(vec[-i])) {  
      return(FALSE)  
    }  
  }  
  return(TRUE)  
}
```

apply

- ▶ Loops in R are inefficient, and best avoided if possible
- ▶ Vectorize operations whenever possible.
- ▶ Using `apply` to *loop* over rows/columns of matrix/table is considered best practice (in terms of clarity)

```
apply(DATA, MARGIN, FUNCTION)
# MARGIN = 1 applies FUNCTION to each row of DATA
# MARGIN = 2 applies FUNCTION to each column of DATA
# MARGIN = c(1,2) applies FUNCTION to both
#     rows and columns of DATA
```

- ▶ e.g., to get the variance of each row/column of a matrix X ,

```
apply(X, 1, var) # variance of rows
apply(X, 2, var) # variance of columns
```

Exercise

The Question

You are given three sticks, each of a random length between 0 and 1.

What's the probability you can make a triangle?

- ▶ The answer is $1/2$
- ▶ Use R to simulate 100,000 times and estimate the answer by
 1. generate 100,000 triplets of uniform (0, 1) random variables
 2. find the portion that can be made into a triangle (hint: use the `is_good` function)

Answer 1: A (not too good) Way

```
N <- 1e5;
ptm <- proc.time() # measure execution time
m <- 0
for (i in 1:N) {
  X <- runif(3)
  if (is_good(X)) {
    m <- m + 1
  }
}
m / N
```

```
## [1] 0.50183
```

```
proc.time() - ptm
```

```
##      user  system elapsed
##    1.187    0.000    1.189
```

Answer 2: An Okay Way

```
N <- 1e5;
ptm <- proc.time() # measure execution time
X <- matrix(runif(N*3), nrow=3, ncol=N)
m <- 0
for (i in 1:N) {
  if (is_good(X[,i])) {
    m <- m + 1
  }
}
m / N
```

```
## [1] 0.50079
```

```
proc.time() - ptm
```

```
##      user  system elapsed
##    0.883    0.000    0.885
```

Answer 3: A Better Way

```
N <- 1e5;  
ptm <- proc.time() # measure execution time  
X <- matrix(runif(N*3), nrow=3, ncol=N)  
m <- apply(X, 2, is_good)  
sum(m) / N
```

```
## [1] 0.50158
```

```
proc.time() - ptm
```

```
##      user  system elapsed  
##    0.793    0.008    0.801
```

Packages

Installing R Packages

- ▶ R has many (*MANY*) packages created by other users that implement state-of-the-art tools (e.g., data manipulation, statistical models)
- ▶ These packages can be downloaded from the Comprehensive R Archive Network (CRAN)
- ▶ This is as simple as running a single line of code:

```
install.packages("package name")
```

- ▶ You will have to select one of many CRAN mirrors (copies across different servers) from which to download the package from
- ▶ For example, to install the `dplyr` package, run

```
install.packages("dplyr")
```

- ▶ You only need to do this *once* for each machine

Loading Packages

- ▶ Once you've installed a package on a machine, you can load the package into your current workspace with the `library()` command
- ▶ For example, to use the `dplyr` package, first load it with

```
library("dplyr")
```