

Session 4 - Statistical Modeling

Jongbin Jung

January 10, 2016

Dependencies

- ▶ Latest version ($\geq 3.1.2$) of R
(*free* from <https://www.r-project.org/>)
- ▶ Latest version of Rstudio (also *free* from <https://www.rstudio.com/>)
- ▶ Packages

```
install.packages('glmnet')  
install.packages('ROCR')
```

Statistical Modeling

- ▶ In this session, I'll demonstrate some common statistical modeling techniques used in R
- ▶ But I *won't*
 - ▶ Go into the theory/details of each (or any) model/method; I assume you know all that already
 - ▶ Introduce the `carat` package, which is a **great** package, but I'm not dealing with it because:
 1. it's pretty well documented here (<http://topepo.github.io/caret/index.html>), and
 2. I (very personally) don't think 'teaching people how to use the `carat` package' is equivalent to 'teaching people statistical modeling in R'
- ▶ I'll focus on regularized linear models with the `glmnet` package, but many of the general concepts work for most models/packages
- ▶ At the very least, you'll end up with a set of (hopefully useful) R snippets

Data Splitting

Pre-processing

```
model.matrix()
```

scale()

Training models

Prediction/Evaluation

Bootstrap